

# Lab 5: Process Loading and Management

Xinyuan (Allen) Pan & Qinyu (Jack) Diao

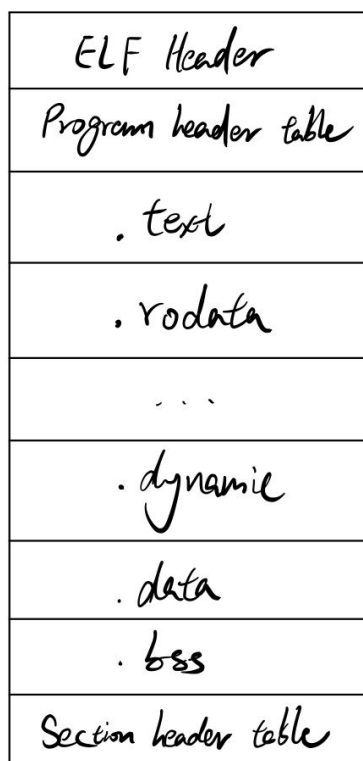
## Objective

The goal of this lab is to add facilities to the OS to load a separately compiled user program from disk, dynamically create and launch an associated OS process, enable the user program to make calls to OS routines, and manage processes and system memory such that a process is removed from the system and all memory occupied by it is reclaimed when its last thread exits.

## Software Design

1) Pictures illustrating the loader operation, showing:

- ELF file layout of compiled user program on disk



- Heap allocation scheme

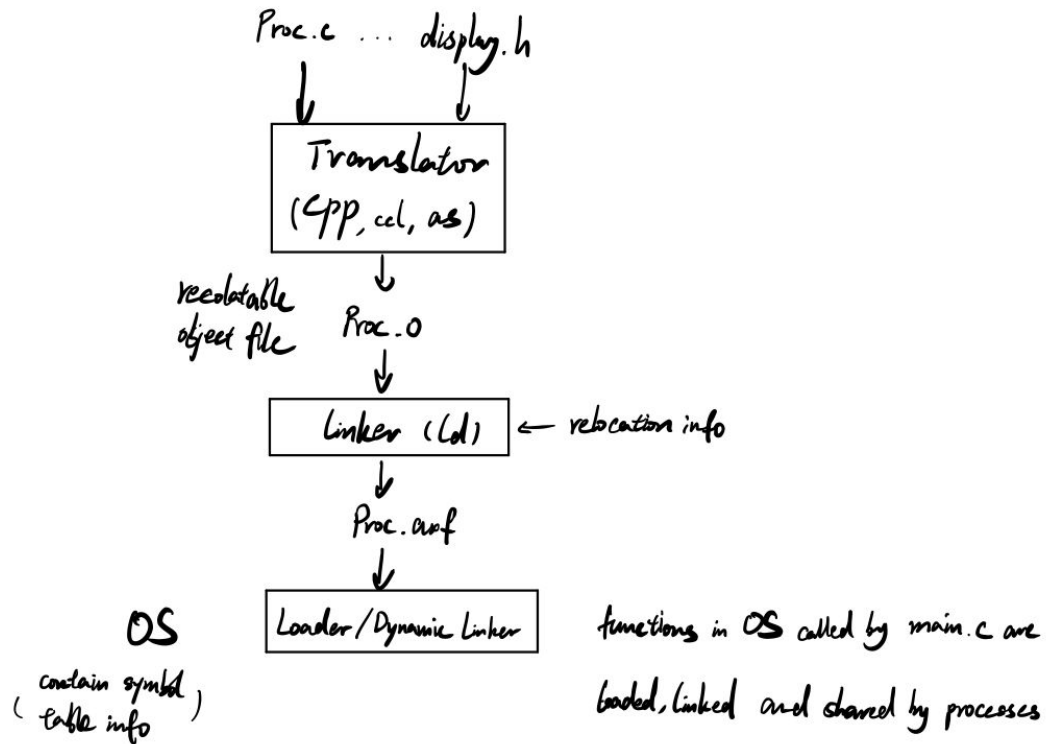
32 B
16 B
64 B
32 B
128 B
64 B
32 B

Free list segregated by size of  $2^h$   
always allocate  $2^k$  bytes

- Memory layout of machine after loading the program

Kernel virtual memory
stack
Memory Mapping Segment
Heap
BSS segment
Data segment
Text segment

- Dynamic linking and relocation process



## 2) Operating system extensions (C and assembly), including SVC\_Handler

**SVC\_Handler:** .func

// Saves R0-R3,R12,LR,PC,PSR (top to bottom)

PUSH {R0,LR}

BL LED\_RED\_TOGGLE

BL LED\_RED\_TOGGLE

POP {R0,LR}

LDR R12,[SP,#24]

// load PC to R12 (the instruction after SVC in user prog)

LDRH R12,[R12,#-2]

// load mem[PC-2] (the SVC instruction) to R12; SVC

instruction is 2 byte

BIC R12,#0xFF00

// &~ clear top bits (extract only the SVC ID)

LDM SP,{R0-R3}

// load any parameters from stack into R0-R3 (auto increment

SP?)

// call corresponding OS\_XXX

PUSH {R4,LR}

// cannot use R0 here, cuz R0 used as return val

CMP R12, #0

BEQ JOS\_Id

CMP R12, #1

BEQ JOS\_Kill

CMP R12, #2

BEQ JOS\_Sleep

CMP R12, #3

BEQ JOS\_Time

CMP R12, #4

BEQ JOS\_AddThread

```

JOS_Id:      BL    OS_Id
              B    done
JOS_Kill:    BL OS_Kill
              B    done
JOS_Sleep:   BL OS_Sleep
              B    done
JOS_Time:    BL OS_Time
              B    done
JOS_AddThread: BL OS_AddThread
              B    done

done:
    POP      {R4,LR}
    STR R0,[SP]                // store return value
    BX LR                    // restore R0-R3,R12,LR,PC,PSR

```

```

int OS_AddProcess(void(*entry)(void), void *text, void *data, unsigned long stackSize, unsigned long priority) {
    static uint32_t nextID = 0;
    unsigned long sr = StartCritical();
    pcbType *newPcb = Heap_Malloc(sizeof(pcbType));
    if (!newPcb) {
        EndCritical(sr);
        return 0;
    }

    newPcb->pid = nextID;
    if (text) newPcb->text = text;                // for OS created processes that do not have text and data on
heap
    else newPcb->text = Heap_Malloc(4);
    if (data) newPcb->data = data;
    else newPcb->data = Heap_Malloc(4);

    newPcb->threadNum = 0;
    dataPt = newPcb->data;
    pcbType *temp = pcbPt; // when adding new process, the thread needs to be added to the new process not
current running process
    pcbPt = newPcb;        // therefore create a temporary pcb to store the current running process
    int in = OS_AddThread(entry, stackSize, priority); // what's the priority here?
    if (!in) {
        killProcess(pcbPt);
        pcbPt = temp;
        EndCritical(sr);
        return 0;
    }

    pcbPt = temp;
    nextID++;
    EndCritical(sr);
    return 1;
}

```

```

// inside setInitialStack
Stacks[i][STACKSIZE-11] = (int32_t) dataPt; // R9

/*
 * Thread Control Block structure
 */
typedef struct tcb {
    int32_t *sp; // ** MUST be the first field ** saved stack pointer (not used by active thread)
    struct tcb *next; // ** MUST be the second field
    struct tcb *prev;
    enum State state;
    int tid;
    uint32_t sleepTimeLeft; // number of cycles left the thread needs to remain in sleep state
    Sema4Type *blocked; // the semaphore it is blocked on
    int32_t priority;
    pcbType *pcb;
} tcbType;

/*
 * Process Control Block Structure
 */
struct pcb {
    int pid;
    void *text;
    void *data;
    int threadNum;
};

```

### 3) High level software system (the new interpreter commands)


```

static void parse_load(char cmd[][20], int len) {
    if (len < 2) {
        Serial_printf("load: missing argument.\n\r");
        return;
    }
    // unsigned long startTime = OS_Time();
    uint32_t startTime = OS_MsTime();
    if (exec_elf(cmd[1], &env) != 1) {
        Serial_printf("load: exec_elf error.\n\r");
        return;
    }
    // unsigned long t = OS_TimeDifference(startTime, OS_Time());
    uint32_t t = OS_MsTime() - startTime;
    Serial_printf("load: process creation time is %u ms\r\n", t);
}

```

# Measurement Data


## 1) Logic analyzer profile of idle task execution




## 2) Logic analyzer profile of user program execution

PF1: SVC traps

PF2 and PF3: LEDs by the user program's main and child threads and process completion. (The duration between the first PF2 toggle and the final PF1 toggle (OS\_Kill) is the time for process completion.)



PF1: process creation time



## Analysis and Discussion

1) Briefly explain the dynamic memory allocation algorithm in your heap manager. Does this implementation have internal or external fragmentation?

The dynamic memory allocation in the heap manager is Knuth's Buddy allocation algorithm. It segregates heap memory into blocks of sizes of  $2^k$  bytes. Whenever requested, it allocates the smallest  $2^k$  number that's larger than the requested size. This way, there is only internal fragmentation and no external fragmentation.

2) How many simultaneously active processes can your system support? What factors limit this number, and how could it be increased?

Our OS can support up to 15 simultaneous active processes (one thread for each). The number of processes (or threads) is limited by the available memory, since stack has fixed size in our design. To increase the number, we can make stack growable.

We can also implement swapping mechanism, so that the data and text segments of some inactive processes can be stored in disk, and only be brought into memory whenever needed.